

R_Deep_Learning_Boston_Housing

Ulises Jose Bustamante Mora

2023-11-26

Importing the Dataset

```
installed.packages("MASS")
```

```
library(MASS)
```

```
head(Boston)
```

```
##      crim zn indus chas   nox   rm  age   dis rad tax ptratio  black
lstat
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900    1 296    15.3 396.90
4.98
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671    2 242    17.8 396.90
9.14
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671    2 242    17.8 392.83
4.03
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622    3 222    18.7 394.63
2.94
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622    3 222    18.7 396.90
5.33
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622    3 222    18.7 394.12
5.21
##      medv
## 1 24.0
## 2 21.6
## 3 34.7
## 4 33.4
## 5 36.2
## 6 28.7
```

Getting more information about the dataset

```
str(Boston)
```

```
## 'data.frame':   506 obs. of  14 variables:
## $ crim   : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn     : num  18 0 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus  : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas   : int   0 0 0 0 0 0 0 0 0 0 ...
## $ nox    : num  0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524
0.524 ...
## $ rm     : num  6.58 6.42 7.18 7 7.15 ...
```

```
## $ age      : num  65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis      : num  4.09 4.97 4.97 6.06 6.06 ...
## $ rad      : int   1 2 2 3 3 3 5 5 5 5 ...
## $ tax      : num  296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num  15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ black    : num  397 397 393 395 397 ...
## $ lstat    : num   4.98 9.14 4.03 2.94 5.33 ...
## $ medv     : num   24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

Checking for n/a values

```
any(is.na(Boston))
```

```
## [1] FALSE
```

```
df = Boston
```

```
head(df)
```

```
##      crim zn indus chas   nox   rm  age    dis rad tax ptratio  black
lstat
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900    1 296    15.3 396.90
4.98
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671    2 242    17.8 396.90
9.14
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671    2 242    17.8 392.83
4.03
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622    3 222    18.7 394.63
2.94
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622    3 222    18.7 396.90
5.33
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622    3 222    18.7 394.12
5.21
##      medv
## 1 24.0
## 2 21.6
## 3 34.7
## 4 33.4
## 5 36.2
## 6 28.7
```

Getting the maximum value of each column

```
maxValue = apply(df,2,max)
maxValue
```

```
##      crim      zn      indus      chas      nox      rm      age
dis
##  88.9762 100.0000  27.7400   1.0000   0.8710   8.7800 100.0000
12.1265
##      rad      tax  ptratio      black      lstat      medv
##  24.0000 711.0000  22.0000 396.9000  37.9700  50.0000
```

Getting the minimum value of each column

```
minValue = apply(df,2,min)
minValue
```

```
##      crim      zn      indus      chas      nox      rm      age
dis
##   0.00632   0.00000   0.46000   0.00000   0.38500   3.56100   2.90000
1.12960
##      rad      tax  ptratio      black      lstat      medv
##   1.00000 187.00000 12.60000   0.32000   1.73000   5.00000
```

Scaling the dataset for better results

```
scaled_df = scale(df, center = minValue, scale = maxValue - minValue)
scaled_df = as.data.frame(scaled_df)
head(scaled_df)
```

```
##      crim  zn      indus chas      nox      rm      age
dis
## 1 0.0000000000 0.18 0.06781525   0 0.3148148 0.5775053 0.6416066
0.2692031
## 2 0.0002359225 0.00 0.24230205   0 0.1728395 0.5479977 0.7826982
0.3489620
## 3 0.0002356977 0.00 0.24230205   0 0.1728395 0.6943859 0.5993821
0.3489620
## 4 0.0002927957 0.00 0.06304985   0 0.1502058 0.6585553 0.4418126
0.4485446
## 5 0.0007050701 0.00 0.06304985   0 0.1502058 0.6871048 0.5283213
0.4485446
## 6 0.0002644715 0.00 0.06304985   0 0.1502058 0.5497222 0.5746653
0.4485446
##      rad      tax  ptratio      black      lstat      medv
## 1 0.00000000 0.20801527 0.2872340 1.0000000 0.08967991 0.4222222
## 2 0.04347826 0.10496183 0.5531915 1.0000000 0.20447020 0.3688889
## 3 0.04347826 0.10496183 0.5531915 0.9897373 0.06346578 0.6600000
## 4 0.08695652 0.06679389 0.6489362 0.9942761 0.03338852 0.6311111
```

```
## 5 0.08695652 0.06679389 0.6489362 1.0000000 0.09933775 0.6933333
## 6 0.08695652 0.06679389 0.6489362 0.9929901 0.09602649 0.5266667
```

Splitting the dataset into train and test data

```
library(caTools)
split = sample.split(scaled_df$medv, SplitRatio = 0.7)
train_df = subset(scaled_df, split == T)
test_df = subset(scaled_df, split == F)

dim(train_df)

## [1] 367 14

dim(test_df)

## [1] 139 14
```

Building the Deep Learning model

```
library(neuralnet)
```

Getting the necessary columns for the model

```
n = names(train_df)

f = as.formula(paste("medv ~", paste(n[!n %in% "medv"], collapse = " +
")))

f

## medv ~ crim + zn + indus + chas + nox + rm + age + dis + rad +
##      tax + ptratio + black + lstat
```

Training the model

```
nn = neuralnet(f, data = train_df, hidden = c(5,3), linear.output = T)
```

Plotting the model

```
plot(nn)
```

Making predictions

```
nn_pre = compute(nn, test_df[1:13])
```

Returning into no scale number for better understanding

```
nn_pre_no_scaled_df = (nn_pre$net.result * (max(df$medv) -  
min(df$medv)) + min(df$medv))
```

```
test_no_scaled_df = (test_df$medv) * (max(df$medv) -  
min(df$medv)) + min(df$medv)
```

Getting the MSE

```
MSE_nn = sum((test_no_scaled_df - nn_pre_no_scaled_df)^2)/nrow(test_df)  
MSE_nn
```

```
## [1] 17.03257
```

Comparing the real value againts the predicted value

```
error_df = data.frame(test_no_scaled_df, nn_pre_no_scaled_df)  
head(error_df)
```

```
##      test_no_scaled_df nn_pre_no_scaled_df  
## 2                21.6             21.65304  
## 7                22.9             20.92027  
## 9                16.5             15.18776  
## 16               19.9             19.86418  
## 18               17.5             18.34040  
## 22               19.6             18.26250
```

Plotting with scatter plot the performance of our model

```
library(ggplot2)
ggplot(error_df, aes(x = test_no_scaled_df, y = nn_pre_no_scaled_df)) +
  geom_point() + stat_smooth()

## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

